

```

import pandas as pd
import re
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_selection import mutual_info_regression, RFE
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler, FunctionTransformer
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.svm import SVR
from xgboost import XGBRegressor

```

```

dataset = pd.read_csv("insurance.csv")
dataset

```

```

{"type": "dataframe", "variable_name": "dataset"}

```

```

dataset.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 58592 entries, 0 to 58591
Data columns (total 45 columns):

```

#	Column	Non-Null Count	Dtype
0	policy_id	58592 non-null	object
1	policy_age	58592 non-null	int64
2	age_of_car	58592 non-null	int64
3	age_of_policyholder	58592 non-null	int64
4	area	58592 non-null	object
5	population_density	58592 non-null	int64
6	make	58592 non-null	object
7	segment	58592 non-null	object
8	model	58592 non-null	object
9	fuel_type	58592 non-null	object
10	max_torque	58592 non-null	object
11	max_power	58592 non-null	object
12	engine_type	58592 non-null	object
13	airbags	58592 non-null	int64
14	is_esc	58592 non-null	object
15	is_adjustable_steering	58592 non-null	object
16	is_tpms	58592 non-null	object
17	is_parking_sensors	58592 non-null	object
18	is_parking_camera	58592 non-null	object
19	rear_brakes_type	58592 non-null	object
20	displacement	58592 non-null	object

21	cylinder	58592	non-null	int64
22	transmission_type	58592	non-null	object
23	gear_box	58592	non-null	int64
24	steering_type	58592	non-null	object
25	turning_radius	58592	non-null	float64
26	length	58592	non-null	object
27	width	58592	non-null	object
28	height	58592	non-null	object
29	gross_weight	58592	non-null	object
30	is_front_fog_lights	58592	non-null	object
31	is_rear_window_wiper	58592	non-null	object
32	is_rear_window_washer	58592	non-null	object
33	is_rear_window_defogger	58592	non-null	object
34	is_brake_assist	58592	non-null	object
35	is_power_door_locks	58592	non-null	object
36	is_central_locking	58592	non-null	object
37	is_power_steering	58592	non-null	object
38	is_driver_seat_height_adjustable	58592	non-null	object
39	is_day_night_rear_view_mirror	58592	non-null	object
40	is_ecw	58592	non-null	object
41	is_speed_alert	58592	non-null	object
42	ncap_rating	58592	non-null	int64
43	claims_in_5_years	58592	non-null	int64
44	claim	58592	non-null	object

dtypes: float64(1), int64(9), object(35)

memory usage: 20.1+ MB

```
def clean_currency_column(column):
    return column.astype(str).apply(lambda x: round(float(re.sub(r'^\d.']), '', x)), 2) if pd.notna(x) and x.strip() else 0)
```

```
dataset['claim'] = clean_currency_column(dataset['claim'])
```

```
dataset['claim'][:5]
```

```
0    773.69
```

```
1   2975.25
```

```
2    682.50
```

```
3   2544.25
```

```
4   3804.65
```

Name: claim, dtype: float64

```
def extract_number_till_letter(text, stop_letter):
    match = re.match(rf'(\d+\.\d*)[{stop_letter}]', text)
    return float(match.group(1)) if match else None
```

```
dataset['max_torque'] = dataset['max_torque'].apply(lambda x:
    extract_number_till_letter(x, "N"))
```

```
dataset['max_power'] = dataset['max_power'].apply(lambda x:
    extract_number_till_letter(x, "b"))
```

```

dataset['length'] = dataset['length'].apply(lambda x:
extract_number_till_letter(x, "c"))
dataset['width'] = dataset['width'].apply(lambda x:
extract_number_till_letter(x, "c"))
dataset['height'] = dataset['height'].apply(lambda x:
extract_number_till_letter(x, "c"))

dataset['gross_weight'] = dataset['gross_weight'].apply(lambda x:
extract_number_till_letter(x, "k"))

dataset['displacement'] = dataset['displacement'].apply(lambda x:
extract_number_till_letter(x, "c"))

dataset.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 58592 entries, 0 to 58591
Data columns (total 45 columns):

```

#	Column	Non-Null Count	Dtype
0	policy_id	58592 non-null	object
1	policy_age	58592 non-null	int64
2	age_of_car	58592 non-null	int64
3	age_of_policyholder	58592 non-null	int64
4	area	58592 non-null	object
5	population_density	58592 non-null	int64
6	make	58592 non-null	object
7	segment	58592 non-null	object
8	model	58592 non-null	object
9	fuel_type	58592 non-null	object
10	max_torque	58592 non-null	float64
11	max_power	58592 non-null	float64
12	engine_type	58592 non-null	object
13	airbags	58592 non-null	int64
14	is_esc	58592 non-null	object
15	is_adjustable_steering	58592 non-null	object
16	is_tpms	58592 non-null	object
17	is_parking_sensors	58592 non-null	object
18	is_parking_camera	58592 non-null	object
19	rear_brakes_type	58592 non-null	object
20	displacement	58592 non-null	float64
21	cylinder	58592 non-null	int64
22	transmission_type	58592 non-null	object
23	gear_box	58592 non-null	int64
24	steering_type	58592 non-null	object
25	turning_radius	58592 non-null	float64
26	length	58592 non-null	float64
27	width	58592 non-null	float64
28	height	58592 non-null	float64
29	gross_weight	58592 non-null	float64

```

30  is_front_fog_lights      58592 non-null object
31  is_rear_window_wiper    58592 non-null object
32  is_rear_window_washer   58592 non-null object
33  is_rear_window_defogger  58592 non-null object
34  is_brake_assist          58592 non-null object
35  is_power_door_locks     58592 non-null object
36  is_central_locking       58592 non-null object
37  is_power_steering        58592 non-null object
38  is_driver_seat_height_adjustable  58592 non-null object
39  is_day_night_rear_view_mirror  58592 non-null object
40  is_ecw                   58592 non-null object
41  is_speed_alert           58592 non-null object
42  ncap_rating              58592 non-null int64
43  claims_in_5_years        58592 non-null int64
44  claim                    58592 non-null float64
dtypes: float64(9), int64(9), object(27)
memory usage: 20.1+ MB

yes_no_columns = ['is_esc', 'is_adjustable_steering', 'is_tpms',
                  'is_parking_sensors', 'is_parking_camera',
                  'is_front_fog_lights', 'is_rear_window_wiper',
                  'is_rear_window_washer', 'is_rear_window_defogger',
                  'is_brake_assist', 'is_power_door_locks',
                  'is_central_locking', 'is_power_steering',
                  'is_driver_seat_height_adjustable',
                  'is_day_night_rear_view_mirror', 'is_ecw', 'is_speed_alert']

onehot_columns = ['area', 'make', 'segment', 'model', 'fuel_type',
                  'engine_type', 'rear_brakes_type',
                  'transmission_type', 'steering_type']

for col in yes_no_columns:
    dataset[col] = dataset[col].map({'Yes':1, 'No':0})

dataset = pd.get_dummies(dataset, columns=onehot_columns,
drop_first=True)

dataset.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 58592 entries, 0 to 58591
Data columns (total 72 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   policy_id                             58592 non-null  object
1   policy_age                             58592 non-null  int64
2   age_of_car                             58592 non-null  int64
3   age_of_policyholder                    58592 non-null  int64
4   area                                   58592 non-null  object
5   population_density                     58592 non-null  int64

```

6	max_torque	58592	non-null	float64
7	max_power	58592	non-null	float64
8	airbags	58592	non-null	int64
9	is_esc	58592	non-null	int64
10	is_adjustable_steering	58592	non-null	int64
11	is_tpms	58592	non-null	int64
12	is_parking_sensors	58592	non-null	int64
13	is_parking_camera	58592	non-null	int64
14	displacement	58592	non-null	float64
15	cylinder	58592	non-null	int64
16	gear_box	58592	non-null	int64
17	turning_radius	58592	non-null	float64
18	length	58592	non-null	float64
19	width	58592	non-null	float64
20	height	58592	non-null	float64
21	gross_weight	58592	non-null	float64
22	is_front_fog_lights	58592	non-null	int64
23	is_rear_window_wiper	58592	non-null	int64
24	is_rear_window_washer	58592	non-null	int64
25	is_rear_window_defogger	58592	non-null	int64
26	is_brake_assist	58592	non-null	int64
27	is_power_door_locks	58592	non-null	int64
28	is_central_locking	58592	non-null	int64
29	is_power_steering	58592	non-null	int64
30	is_driver_seat_height_adjustable	58592	non-null	int64
31	is_day_night_rear_view_mirror	58592	non-null	int64
32	is_ecw	58592	non-null	int64
33	is_speed_alert	58592	non-null	int64
34	ncap_rating	58592	non-null	int64
35	claims_in_5_years	58592	non-null	int64
36	claim	58592	non-null	float64
37	make_Ford	58592	non-null	bool
38	make_GM	58592	non-null	bool
39	make_Honda	58592	non-null	bool
40	make_Toyota	58592	non-null	bool
41	segment_B1	58592	non-null	bool
42	segment_B2	58592	non-null	bool
43	segment_C1	58592	non-null	bool
44	segment_C2	58592	non-null	bool
45	segment_Utility	58592	non-null	bool
46	model_M10	58592	non-null	bool
47	model_M11	58592	non-null	bool
48	model_M2	58592	non-null	bool
49	model_M3	58592	non-null	bool
50	model_M4	58592	non-null	bool
51	model_M5	58592	non-null	bool
52	model_M6	58592	non-null	bool
53	model_M7	58592	non-null	bool
54	model_M8	58592	non-null	bool

```

55 model_M9 58592 non-null bool
56 fuel_type_Diesel 58592 non-null bool
57 fuel_type_Petrol 58592 non-null bool
58 engine_type_1.2 L K Series Engine 58592 non-null bool
59 engine_type_1.2 L K12N Dualjet 58592 non-null bool
60 engine_type_1.5 L U2 CRDi 58592 non-null bool
61 engine_type_1.5 Turbocharged Revotorq 58592 non-null bool
62 engine_type_1.5 Turbocharged Revotron 58592 non-null bool
63 engine_type_F8D Petrol Engine 58592 non-null bool
64 engine_type_G12B 58592 non-null bool
65 engine_type_K Series Dual jet 58592 non-null bool
66 engine_type_K10C 58592 non-null bool
67 engine_type_i-DTEC 58592 non-null bool
68 rear_brakes_type_Drum 58592 non-null bool
69 transmission_type_Manual 58592 non-null bool
70 steering_type_Manual 58592 non-null bool
71 steering_type_Power 58592 non-null bool
dtypes: bool(35), float64(9), int64(26), object(2)
memory usage: 18.5+ MB

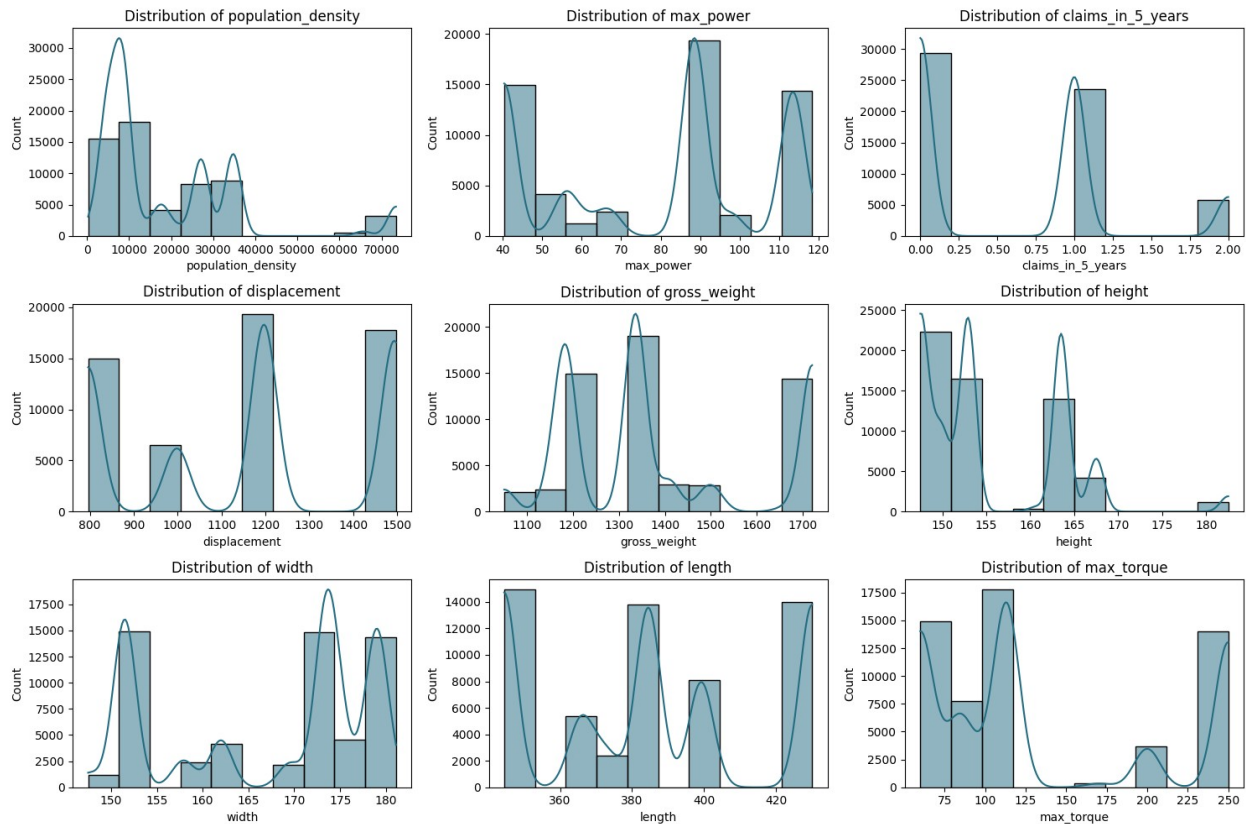
continuous_features = ['population_density', 'max_power',
                        'claims_in_5_years', 'displacement',
                        'gross_weight', 'height', 'width', 'length',
                        'max_torque']

plt.figure(figsize=(15, 10))

for i, col in enumerate(continuous_features, 1):
    plt.subplot(3, 3, i)
    sns.histplot(dataset[col], kde=True, bins=10, color='#246D82')
    plt.title(f"Distribution of {col}")

plt.tight_layout()
plt.show()

```



```
log_transformer = FunctionTransformer(np.log1p, validate=True)
scaler = StandardScaler()
```

```
dataset['max_power'] =
log_transformer.fit_transform(dataset[['max_power']])
dataset['displacement'] =
log_transformer.fit_transform(dataset[['displacement']])
```

```
for col in ['gross_weight', 'height', 'width', 'length',
'max_torque']:
    dataset[col] = scaler.fit_transform(dataset[[col]])
```

```
dataset.head()
```

```
{"type": "dataframe", "variable_name": "dataset"}
```

```
dataset2=dataset.copy()
```

```
dataset = dataset.drop(columns=['policy_id'])
```

```
X = dataset.drop(columns=['claim', 'population_density', 'area'])
y = dataset['claim']
```

```
corr_matrix =
dataset.drop(columns=['population_density', 'area']).corr()
['claim'].abs().sort_values(ascending=False)[1:11]
```

```

rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X, y)
rf_importances = pd.Series(rf.feature_importances_,
index=X.columns).sort_values(ascending=False).head(10)

mi_importances = pd.Series(mutual_info_regression(X, y,
random_state=42),
index=X.columns).sort_values(ascending=False).head(10)

rfe = RFE(estimator=LinearRegression(), n_features_to_select=10)
rfe.fit(X, y)
rfe_importances = pd.Series(rfe.support_, index=X.columns)
rfe_selected_features = rfe_importances[rfe_importances == True].index

fig, axes = plt.subplots(2, 2, figsize=(14, 10))
fig.suptitle('Feature Importance Across Different Methods')

sns.barplot(x=corr_matrix.values, y=corr_matrix.index, ax=axes[0, 0],
color='blue')
axes[0, 0].set_title('Correlation with Claim')

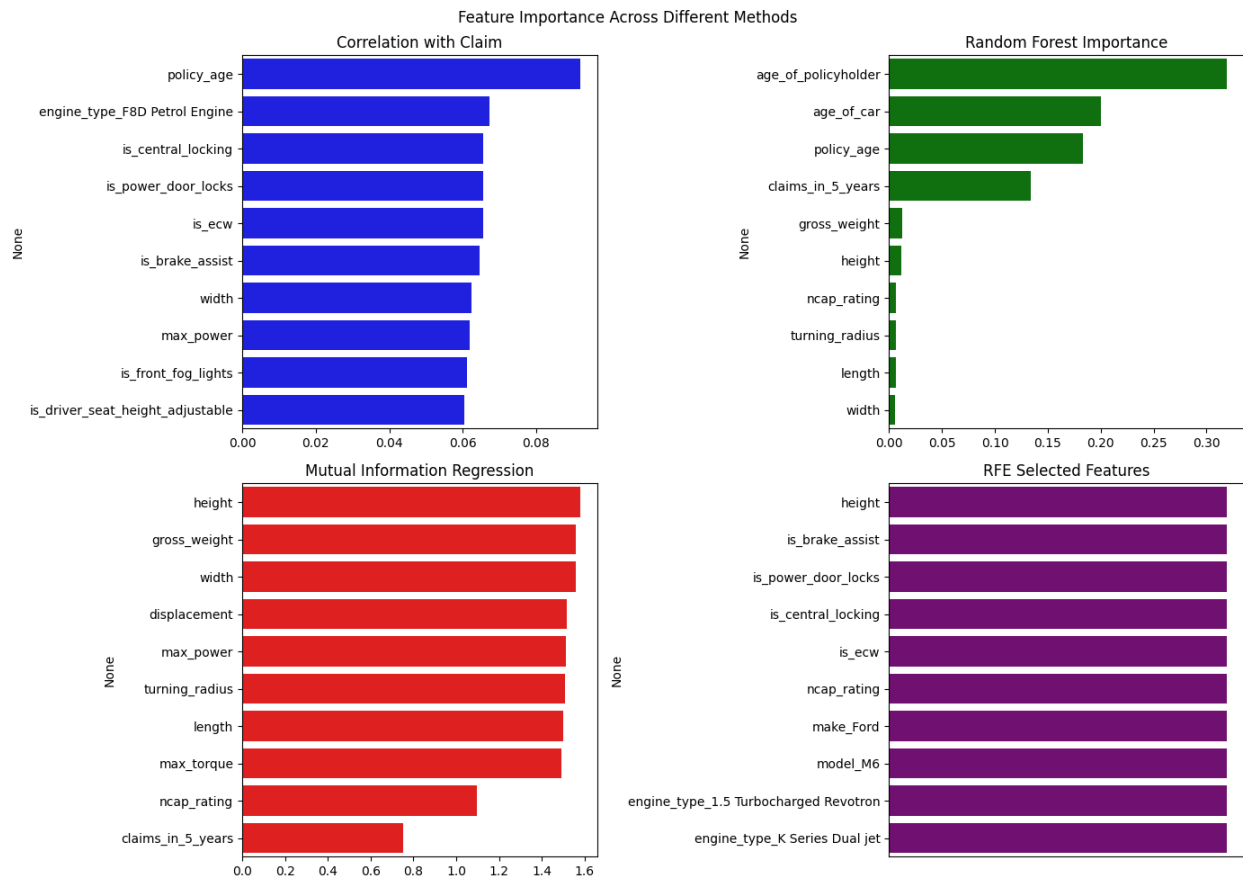
sns.barplot(x=rf_importances.values, y=rf_importances.index,
ax=axes[0, 1], color='green')
axes[0, 1].set_title('Random Forest Importance')

sns.barplot(x=mi_importances.values, y=mi_importances.index,
ax=axes[1, 0], color='red')
axes[1, 0].set_title('Mutual Information Regression')

sns.barplot(x=[1] * len(rfe_selected_features),
y=rfe_selected_features, ax=axes[1, 1], color='purple')
axes[1, 1].set_title('RFE Selected Features')
axes[1, 1].set_xticks([])

plt.tight_layout()
plt.show()

```

mi_importances

```
height 1.578313
gross_weight 1.558201
width 1.556319
displacement 1.514891
max_power 1.510224
turning_radius 1.508697
length 1.497975
max_torque 1.492158
ncap_rating 1.097225
claims_in_5_years 0.750461
dtype: float64
```

rf_importances

```
age_of_policyholder 0.319349
age_of_car 0.200270
policy_age 0.183155
claims_in_5_years 0.134020
gross_weight 0.012498
height 0.011694
ncap_rating 0.006433
```

```


turning_radius      0.006429
length             0.006161
width              0.005678
dtype: float64

corr_matrix =
dataset.drop(columns=['population_density', 'area', 'claim']).corr()

plt.figure(figsize=(12, 8))

sns.heatmap(corr_matrix, fmt=".2f", cmap="coolwarm", linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()

```



```

important_features = ['max_power',
                      'is_ecw',
                      'is_front_fog_lights',
                      'claims_in_5_years',
                      'gross_weight',

```

```

'policy_age',
'ncap_rating',
'is_power_steering',
'height',
'width',
'age_of_policyholder',
'is_brake_assist',
'displacement',
'max_torque',
'is_driver_seat_height_adjustable',
'is_central_locking',
'length',
"area_Boston",
"area_Charlotte",
"area_Chicago",
"area_Columbus",
"area_Dallas",
"area_Denver",
"area_Fort Worth",
"area_Houston",
"area_Indianapolis",
"area_Jacksonville",
"area_Los Angeles",
"area_Nashville",
"area_New York",
"area_Philadelphia",
"area_Phoenix",
"area_San Antonio",
"area_San Diego",
"area_San Francisco",
"area_San Jose",
"area_Seattle",
"area_Washington D.C." ]

```

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 58592 entries, 0 to 58591
```

```
Data columns (total 91 columns):
```

#	Column	Non-Null Count		Dtype
0	policy_age	58592	non-null	int64
1	age_of_car	58592	non-null	int64
2	age_of_policyholder	58592	non-null	int64
3	population_density	58592	non-null	int64
4	max_torque	58592	non-null	float64
5	max_power	58592	non-null	float64
6	airbags	58592	non-null	int64
7	is_esc	58592	non-null	int64
8	is_adjustable_steering	58592	non-null	int64

9	is_tpms	58592	non-null	int64
10	is_parking_sensors	58592	non-null	int64
11	is_parking_camera	58592	non-null	int64
12	displacement	58592	non-null	float64
13	cylinder	58592	non-null	int64
14	gear_box	58592	non-null	int64
15	turning_radius	58592	non-null	float64
16	length	58592	non-null	float64
17	width	58592	non-null	float64
18	height	58592	non-null	float64
19	gross_weight	58592	non-null	float64
20	is_front_fog_lights	58592	non-null	int64
21	is_rear_window_wiper	58592	non-null	int64
22	is_rear_window_washer	58592	non-null	int64
23	is_rear_window_defogger	58592	non-null	int64
24	is_brake_assist	58592	non-null	int64
25	is_power_door_locks	58592	non-null	int64
26	is_central_locking	58592	non-null	int64
27	is_power_steering	58592	non-null	int64
28	is_driver_seat_height_adjustable	58592	non-null	int64
29	is_day_night_rear_view_mirror	58592	non-null	int64
30	is_ecw	58592	non-null	int64
31	is_speed_alert	58592	non-null	int64
32	ncap_rating	58592	non-null	int64
33	claims_in_5_years	58592	non-null	int64
34	claim	58592	non-null	float64
35	area_Boston	58592	non-null	bool
36	area_Charlotte	58592	non-null	bool
37	area_Chicago	58592	non-null	bool
38	area_Columbus	58592	non-null	bool
39	area_Dallas	58592	non-null	bool
40	area_Denver	58592	non-null	bool
41	area_Fort Worth	58592	non-null	bool
42	area_Houston	58592	non-null	bool
43	area_Indianapolis	58592	non-null	bool
44	area_Jacksonville	58592	non-null	bool
45	area_Los Angeles	58592	non-null	bool
46	area_Nashville	58592	non-null	bool
47	area_New York	58592	non-null	bool
48	area_Philadelphia	58592	non-null	bool
49	area_Phoenix	58592	non-null	bool
50	area_San Antonio	58592	non-null	bool
51	area_San Diego	58592	non-null	bool
52	area_San Francisco	58592	non-null	bool
53	area_San Jose	58592	non-null	bool
54	area_Seattle	58592	non-null	bool
55	area_Washington D.C.	58592	non-null	bool
56	make_Ford	58592	non-null	bool
57	make_GM	58592	non-null	bool

58	make_Honda	58592	non-null	bool
59	make_Toyota	58592	non-null	bool
60	segment_B1	58592	non-null	bool
61	segment_B2	58592	non-null	bool
62	segment_C1	58592	non-null	bool
63	segment_C2	58592	non-null	bool
64	segment_Utility	58592	non-null	bool
65	model_M10	58592	non-null	bool
66	model_M11	58592	non-null	bool
67	model_M2	58592	non-null	bool
68	model_M3	58592	non-null	bool
69	model_M4	58592	non-null	bool
70	model_M5	58592	non-null	bool
71	model_M6	58592	non-null	bool
72	model_M7	58592	non-null	bool
73	model_M8	58592	non-null	bool
74	model_M9	58592	non-null	bool
75	fuel_type_Diesel	58592	non-null	bool
76	fuel_type_Petrol	58592	non-null	bool
77	engine_type_1.2 L K Series Engine	58592	non-null	bool
78	engine_type_1.2 L K12N Dualjet	58592	non-null	bool
79	engine_type_1.5 L U2 CRDi	58592	non-null	bool
80	engine_type_1.5 Turbocharged Revotorq	58592	non-null	bool
81	engine_type_1.5 Turbocharged Revotron	58592	non-null	bool
82	engine_type_F8D Petrol Engine	58592	non-null	bool
83	engine_type_G12B	58592	non-null	bool
84	engine_type_K Series Dual jet	58592	non-null	bool
85	engine_type_K10C	58592	non-null	bool
86	engine_type_i-DTEC	58592	non-null	bool
87	rear_brakes_type_Drum	58592	non-null	bool
88	transmission_type_Manual	58592	non-null	bool
89	steering_type_Manual	58592	non-null	bool
90	steering_type_Power	58592	non-null	bool

dtypes: bool(56), float64(9), int64(26)
memory usage: 18.8 MB

```

X = dataset[important_features]
y = dataset['claim']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

models = {
    "Linear Regression": LinearRegression(),
    "Random Forest": RandomForestRegressor(n_estimators=100,
random_state=42),
    "Gradient Boosting": GradientBoostingRegressor(n_estimators=100,
random_state=42),
    "Support Vector Regressor": SVR(),
    "XGBoost": XGBRegressor(n_estimators=100, random_state=42)

```

```

}

results = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    results[name] = {"MSE": mse, "RMSE": rmse, "MAE": mae, "R2 Score":
r2}

results_df = pd.DataFrame(results).T

fig, axes = plt.subplots(2, 2, figsize=(14, 10))

sns.barplot(x=results_df.index, y=results_df['MSE'], ax=axes[0, 0],
color='blue')
axes[0, 0].set_title('Mean Squared Error (MSE)')
axes[0, 0].set_ylabel('MSE Value')
axes[0, 0].set_xticklabels(results_df.index, rotation=20)

sns.barplot(x=results_df.index, y=results_df['RMSE'], ax=axes[0, 1],
color='green')
axes[0, 1].set_title('Root Mean Squared Error (RMSE)')
axes[0, 1].set_ylabel('RMSE Value')
axes[0, 1].set_xticklabels(results_df.index, rotation=20)

sns.barplot(x=results_df.index, y=results_df['MAE'], ax=axes[1, 0],
color='red')
axes[1, 0].set_title('Mean Absolute Error (MAE)')
axes[1, 0].set_ylabel('MAE Value')
axes[1, 0].set_xticklabels(results_df.index, rotation=20)

sns.barplot(x=results_df.index, y=results_df['R2 Score'], ax=axes[1,
1], color='purple')
axes[1, 1].set_title('R2 Score')
axes[1, 1].set_ylabel('R2 Value')
axes[1, 1].set_xticklabels(results_df.index, rotation=20)

plt.tight_layout()
plt.show()

<ipython-input-34-e95f68489882>:7: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
    axes[0, 0].set_xticklabels(results_df.index, rotation=20)

```

```
<ipython-input-34-e95f68489882>:13: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
```

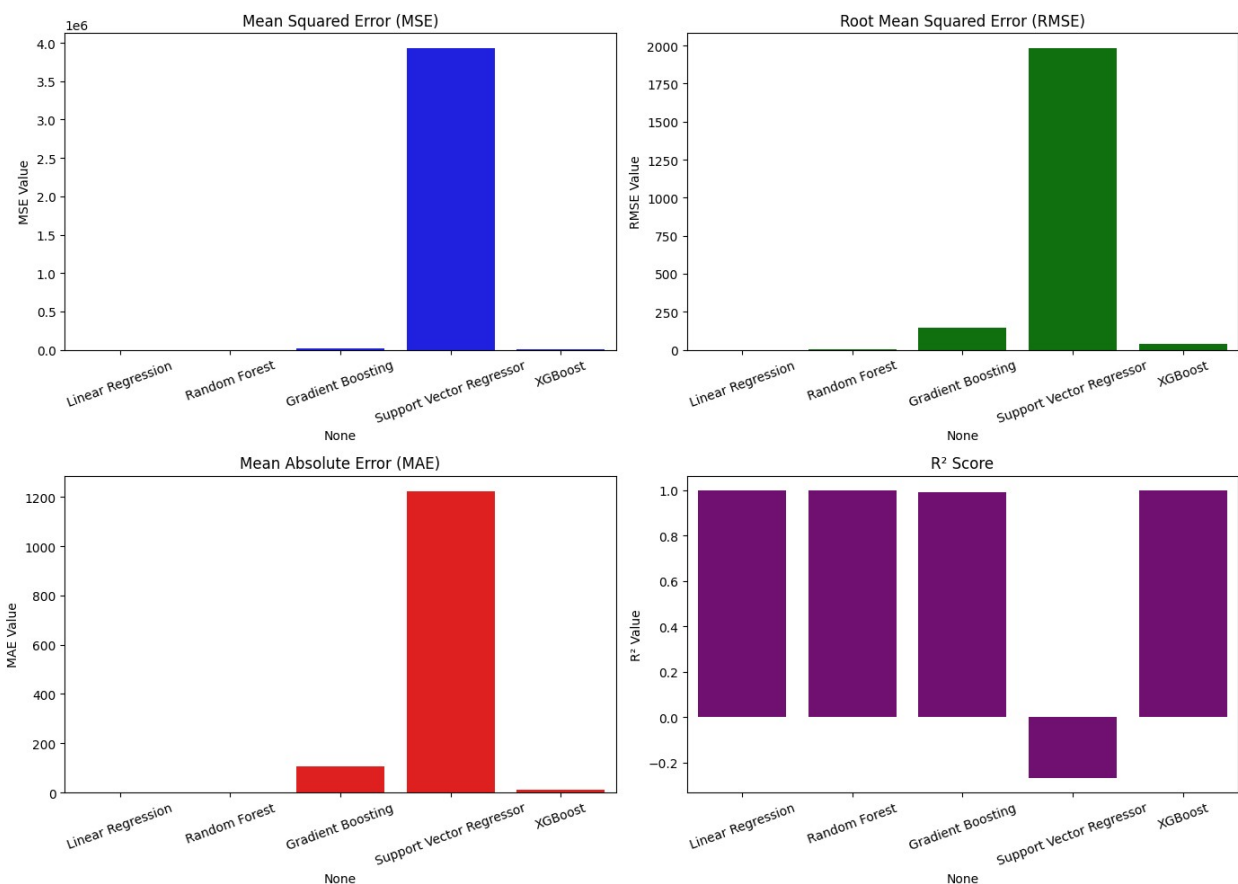
```
axes[0, 1].set_xticklabels(results_df.index, rotation=20)
```

```
<ipython-input-34-e95f68489882>:19: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
```

```
axes[1, 0].set_xticklabels(results_df.index, rotation=20)
```

```
<ipython-input-34-e95f68489882>:25: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
```

```
axes[1, 1].set_xticklabels(results_df.index, rotation=20)
```



```
results_df
```

```
{ "summary": "{\n  \"name\": \"results_df\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"MSE\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1755024.4142618526,\n        \"min\": 0.0692928672673422,\n        \"max\": 3929820.6522851484,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          1.0316953587004944,\n          1379.9056242508052,\n          20682.220268141984\n        ],\n        \"na\": 0\n      }\n    }\n  ]\n}
```

```

{"semantic_type": "", "description": "", "column": "RMSE", "properties": {"dtype": "number", "std": 868.1547643785032, "min": 0.2632353837677264, "max": 1982.3775251664724, "num_unique_values": 5, "samples": [1.0157240563757928, 37.147080965411064, 143.8131435861896]}, "semantic_type": "", "description": "", "column": "MAE", "properties": {"dtype": "number", "std": 536.3139785710244, "min": 0.2070779678762725, "max": 1224.0919539652052, "num_unique_values": 5, "samples": [0.2530196136192616, 10.50190264170198, 104.56140340786972]}, "semantic_type": "", "description": "", "column": "R\u00b2 Score", "properties": {"dtype": "number", "std": 0.5664842536290661, "min": -0.2684618521629818, "max": 0.999999977633748, "num_unique_values": 5, "samples": [0.999999669908575, 0.9995545960493313, 0.9933242227191933]}, "semantic_type": "", "description": "", "column": "type", "type": "dataframe", "variable_name": "results_df"}

```